

Training-window sensitivity with full-history evaluation



Fixed student-disjoint split, DKT-64, AdamW (weight decay 0), batch size 64, eight-epoch cap and full-history test evaluation. The shared 0.75–0.78 scale avoids visually magnifying the small mean differences. Error bars are across-seed SDs; this is a descriptive sensitivity analysis, not a confidence interval or equivalence test.